



N°d'ordre NNT : ?

THÈSE de DOCTORAT DE L'UNIVERSITÉ DE LILLE

École Doctorale MADIS N° 631

**Mathématiques-Sciences du numérique et de leurs
interactions**

Spécialité / discipline de doctorat : Mathématique

Soutenue par :

Maxime Haddouche

**On the Interplays between Generalisation and
Optimisation: a PAC-Bayes Approach**

ACKNOWLEDGEMENTS

TODO

Quote 1

Author

Quote 2

AUTHOR

CONTENTS

Contents	5
List of Theorems	8
List of Notations	9
List of Publications	11
Conference article	11
Journal article	11
Research Report	11
Préambule: Apprentissage Humain, Apprentissage Machine et Généralisation	13
Preamble: Human Learning, Machine Learning and Generalisation	17
1 PAC-Bayes Learning, a field of many paradigms	21
1.1 A brief introduction to statistical learning	21
1.2 An information-theoretic exposition of PAC-Bayes learning	23
1.3 From theory to learning algorithms	27
1.4 An optimisation perspective of PAC-Bayes	30
2 PAC-Bayesian Bounds for Martingales and Heavy-Tailed losses	37
3 Mitigating Initialisation Impact by Real-Time Control: Online PAC- Bayes Learning	39
4 Mitigating Initialisation Impact through Flat Minima: Fast Rates for Small Gradients	41
4.1 Introduction	42
5 Wasserstein PAC-Bayes Learning: Exploiting Optimisation Guarantees to Explain Generalisation	43
5.1 Introduction	44
A Some Mathematical Tools	45
A.1 JENSEN's Inequality	45

A.2	MARKOV's Inequality	45
A.3	Ville's Inequality	46
B	Additional Background	47
B.1	Online Learning	47
B.2	Wasserstein distances and optimisation on distributions spaces	47
B.3	Differential privacy	47
B.4	Log-Sobolev and Poincaré Inequalities	47
	References	49

LIST OF ALGORITHMS

Chapter 1

Chapter 2

Chapter 3

Chapter 5

LIST OF THEOREMS

Chapter 1

1.1.1 Example	22
1.1.1 Definition	23
1.2.1 Proposition	24
Proof (of Equation (1.4))	25
1.2.1 Lemma (Change of measure inequality)	25
1.3.1 Example	27

Chapter 2

Chapter 3

Chapter 4

Chapter 5

Appendix

A.1.1 Theorem (JENSEN's Inequality)	45
A.2.1 Theorem (MARKOV's Inequality)	45
A.3.1 Lemma (Ville's maximal inequality for supermartingales)	46

LIST OF NOTATIONS

General

a	A scalar (integer or real)
\mathbb{R}	The set of real numbers
\mathbb{R}^d	The euclidean set of dimension d
$\ \cdot\ $	a norm of an euclidean set
$\text{dist}(\cdot, \cdot)$	A distance on a Polish space.
\mathbb{N}	The set of natural numbers
∇f	the gradient of a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$

Statistical Learning Theory

\mathcal{Z}	Data space. In supervised learning, $\mathcal{Z} = (\mathcal{X}, \mathcal{Y})$ with \mathcal{X}, \mathcal{Y} input and label spaces
\mathbf{z}	A datum of \mathcal{Z} , in supervised learning $\mathbf{z} = (\mathbf{x}, y)$ with \mathbf{x} input and y label
\mathcal{S}	Learning sample $\mathcal{S} = \{\mathbf{z}_i\}_{i \geq 1}$
$\mathcal{D}_{\mathcal{S}}$	Distribution of \mathcal{S}
\mathcal{S}_m	Restriction of \mathcal{S} to its m first data $\mathcal{S}_m = \{\mathbf{z}_i\}_{i=1 \dots m}$
\mathcal{D}	For <i>i.i.d.</i> \mathcal{S} , distribution of a single datum on \mathcal{Z}
\mathcal{D}^m	For <i>i.i.d.</i> \mathcal{S} , distribution of \mathcal{S}_m
\mathcal{T}	For <i>i.i.d.</i> \mathcal{S} , Test set drawn from \mathcal{D}
\mathcal{H}	The set of hypotheses
h	A hypothesis $h \in \mathcal{H}$
ℓ	Loss function $\ell : \mathcal{H} \times \mathcal{Z} \rightarrow \mathbb{R}$

Probability Theory

$\mathbb{E}_{X \sim \nu} [\cdot]$	The expectation <i>w.r.t.</i> the random variable $X \sim \nu$
$\mathbb{P}_{X \sim \nu} [\cdot]$	The probability <i>w.r.t.</i> the random variable $X \sim \nu$
$\mathbb{1} [a]$	Indicator function; returns 1 if a is true and 0 otherwise
$(\mathcal{F}_i)_{i \geq 1}$	Filtration adapted to \mathcal{S}
$\mathbb{E}_i[\cdot]$	Conditional expectation <i>w.r.t.</i> \mathcal{F}_i , i.e. $\mathbb{E}_i[\cdot] := \mathbb{E}[\cdot \mid \mathcal{F}_i]$
$\mathcal{N}(\mu, \Sigma)$	Gaussian distribution on \mathbb{R}^d with mean μ and covariance matrix Σ

PAC-Bayes framework

$\mathcal{M}(\mathcal{H})$	Set of Probability densities <i>w.r.t.</i> the reference measure on \mathcal{H}
Q	Posterior distribution $Q \in \mathcal{M}(\mathcal{H})$ on \mathcal{H}
P	Prior distribution $P \in \mathcal{M}(\mathcal{H})$ on \mathcal{H}
$KL(Q \ P)$	Kullback-Leibler (KL) divergence between Q and P
$D_\alpha(Q \ P)$	Rényi Divergence between Q and P
$R_{\mathcal{D}}(h)$	Population Risk of $h \in \mathcal{H}$ <i>w.r.t.</i> \mathcal{D} , i.e. $R_{\mathcal{D}}(h) := \mathbb{E}_{\mathbf{z} \sim \mathcal{D}} [\ell(h, \mathbf{z})]$
$\hat{R}_{\mathcal{S}_m}(h)$	Empirical Risk on \mathcal{S}_m , i.e. $\hat{R}_{\mathcal{S}_m}(h) = \frac{1}{m} \sum_{i=1}^m \ell(h, \mathbf{z}_i)$
$\Delta_{\mathcal{S}_m}(h)$	Generalisation gap $\Delta_{\mathcal{S}_m}(h) := R_{\mathcal{D}}(h) - \hat{R}_{\mathcal{S}_m}(h)$
$R_{\mathcal{D}}(Q)$	Expected population risk <i>w.r.t.</i> Q , i.e. $R_{\mathcal{D}}(Q) := \mathbb{E}_{h \sim Q} [R_{\mathcal{D}}(h)]$
$\hat{R}_{\mathcal{S}_m}(Q)$	Expected empirical risk <i>w.r.t.</i> Q , $\hat{R}_{\mathcal{S}_m}(Q) := \mathbb{E}_{h \sim Q} [\hat{R}_{\mathcal{S}_m}(h)]$
$\Delta_{\mathcal{S}_m}(Q)$	Expected generalisation gap <i>w.r.t.</i> Q , $\Delta_{\mathcal{S}_m}(Q) := \mathbb{E}_{h \sim Q} [\Delta_{\mathcal{S}_m}(h)]$
$P_{-f(h)}$	Gibbs posterior associated to prior P and function $f : \mathcal{H} \rightarrow \mathbb{R}$

Optimal transport

W_1	The 1-Wasserstein distance
W_2	The 2-Wasserstein distance
$\Gamma(Q, P)$	Set of all coupling distribution on \mathcal{H}^2 whose marginals are Q and P .

LIST OF PUBLICATIONS

Conference article

PAUL VIALARD, MAXIME HADDOUCHE, Umut SIMSEKLI, and BENJAMIN GUEDJ. Learning via Wasserstein-Based High Probability Generalisation Bounds. *Advances in Neural Information Processing Systems (NeurIPS)*. (2023).

MAXIME HADDOUCHE and BENJAMIN GUEDJ. Online PAC-Bayes Learning. *Advances in Neural Information Processing Systems (NeurIPS)*. (2022).

Journal article

MAXIME HADDOUCHE and BENJAMIN GUEDJ. PAC-Bayes Generalisation Bounds for Heavy-Tailed Losses through Supermartingales. *Transactions on Machine Learning Research*. (2023).

Research Report

MAXIME HADDOUCHE and BENJAMIN GUEDJ. Wasserstein PAC-Bayes Learning: A Bridge Between Generalisation and Optimisation. *arXiv*. abs/2304.07048. (2023).

MAXIME HADDOUCHE, BENJAMIN GUEDJ, and OLIVIER WINTENBERGER. Optimistic Dynamic Regret Bounds. (2023). *arXiv*: 2301.07530 [cs.LG].

PIERRE JOBIC, MAXIME HADDOUCHE, and BENJAMIN GUEDJ. Federated Learning with Nonvacuous Generalisation Bounds. (2023). *arXiv*: 2310.11203 [cs.LG].

MAXIME HADDOUCHE, BENJAMIN GUEDJ, and JOHN SHAWE-TAYLOR. Upper and Lower Bounds on the Performance of Kernel PCA. (2020). *arXiv*: 2012.10369 [cs.LG].

PRÉAMBULE: APPRENTISSAGE HUMAIN, APPRENTISSAGE MACHINE ET GÉNÉRALISATION

Ce manuscrit étudie la question de la capacité de *généralisation* des algorithmes d'apprentissage machine. Pour comprendre la généralisation, il faut d'abord appréhender l'*apprentissage*, prenons donc le luxe, pour un bref instant, d'oublier les machines pour se concentrer sur l'apprentissage en ce qu'il a de plus humain.

Appréhender l'apprentissage humain. Un être apprenant, en premier lieu, va se structurer autour d'expériences, vécues ou transmises par autrui et va ensuite en bénéficier via diverses modalités. Il peut, par exemple, considérer une expérience médiée comme vraie (le feu brûle) et agir en fonction de ce postulat alors qu'à l'opposé, la répétition ou la négation de cette même expérience peuvent être symptomatiques d'une valeur de vérité nulle. Ces scénarios peuvent tout aussi bien apparaître pour une expérience vécue (hallucinations). Cette première dichotomie quant au traitement de l'information est intrinsèquement liée à une question clairement énoncée : est-ce que le feu brûle ? Puis-je me fier à mes sens ou ai-je halluciné ? Dans ces cas de figure, l'apprentissage a eu lieu à travers l'assujettissement de l'expérience à sa valeur de vérité par rapport à une question simple (ici à deux issues). Cette vision peut facilement s'étendre à une arborescence finie de possibles pour des questions à choix multiples. En effet, on peut étendre la question de la brûlure comme suit : quelle est l'intensité de la brûlure en fonction de la température du feu ? On peut dès lors établir une multitude de réponses représentant divers degrés de brûlure.

Néanmoins, de nombreuses questions ne peuvent se réduire à un nombre fini de possibilités. Par exemple, qu'est-ce que le feu ? Pour répondre à cette question, il est néanmoins possible d'exploiter de multiples facettes d'expériences (feu de bois, brindille, roche) pour proposer le feu comme étant la réaction chimique de l'oxygène de l'air avec un matériau combustible, un apport d'énergie servant de déclencheur.

Il est alors légitime de se demander pourquoi l'apprenant a eu besoin de comprendre la vraie nature du feu. Cette compréhension fondamentale des choses émerge de considérations pratiques : comment ne plus avoir froid ? Peut-on manger de la viande autrement que crue pour diminuer les risques de maladie ? Il faut alors de multiples interactions avec l'environnement pour générer des expériences et ensuite apprendre d'elles pour répondre graduellement à un besoin complexe (comment faire un feu pour se réchauffer?).

Ainsi, par cette analyse préliminaire, nous avons trouvé plusieurs prémices de compréhension de l'apprentissage chez l'homme.

- Comment l'apprentissage se formalise-t-il structurellement ? L'apprenant doit abâtardir l'expérience à des questions simples pour acquérir des certitudes primaires. Ces dernières acquises, il est possible d'atteindre des questions complexes en imbriquant de plus en plus de considérations élémentaires.
- D'où provient le besoin d'apprendre ? D'un point de vue pratique, l'émergence de ces questions complexes dérive bien souvent d'un rapport de l'être à son environnement, permettant d'élaborer des objectifs contextuels. L'apprenant devient alors graduellement capable de répondre à des besoins complexes par une succession d'actions simples.

De l'apprentissage humain à l'apprentissage machine. L'apprentissage machine s'est structuré autour de deux approches, une première symbolique qui tire profit des extrapolations humaines pour apprendre à la machine à manipuler une axiomatique et une seconde, statistique, qui consiste à fournir bon nombre d'expériences à la machine pour lui faire apprendre par de multiples exemples empiriques. Nous allons nous focaliser sur la seconde approche car, elle sous-tend une large partie de la recherche moderne. Cette méthode requiert de nombreuses expériences transmises à la machine qui en extrait les connaissances à travers des procédures optimisatoires. Plus précisément, la connaissance extraite dépend de la question posée ainsi que sa traduction mathématique. Nous pouvons alors relever des parallèles avec l'apprentissage humain décrit plus haut: il faut des expériences et une question pour réduire le réel à quelque chose d'apprenable. Pour aller plus loin, la variété des scénarii d'apprentissages humains décrits au dessus ont une correspondance dans l'apprentissage machine moderne: à la question "Le feu brûle-t-il?" on peut associer l'apprentissage supervisé qui traite d'apprendre sur des questions à choix multiples. A la question "qu'est-ce que le feu?", on peut associer l'apprentissage non-supervisé qui va chercher, dans le cas du clustering (ou regroupement), des similitudes non-induites par la question entre diverses expériences. Finalement, quant à l'interaction avec l'environnement et la question "puis-je faire un feu?", elle est associée à l'apprentissage par renforcement qui étudie l'apprentissage d'un agent qui interagit avec son environnement.

Comprendre la généralisation depuis l'apprentissage. La généralisation peut être vue comme la capacité d'exploiter l'apprentissage d'une expérience au delà de cette dernière. Cela englobe une compréhension théorique et axiomatique d'un phénomène bien au delà de l'expérience en elle-même, *i.e.* une extrapolation fructueuse ou bien la capacité à exploiter la connaissance acquise pour une situation inédite, présentant des similitudes avec divers vécus, *i.e.* interpoler des expériences.

Ce double aspect de la généralisation se retrouve aussi bien chez l'homme que la machine sous diverses modalités. Les réseaux de neurones profonds, qui sont le fer de lance de l'apprentissage machine moderne, se basent sur des espaces de dimension finie pour apprendre, ce qui revient à dire qu'un problème peut être appris à travers un nombre fini de principes fondateurs. Le nombre de principes pouvant être augmentés autant que les capacités numériques le permettent, nous dirons alors que les réseaux de neurones ont une puissance discrète de généralisation. Etant donné que les méthodes d'apprentissage machine sont corrélées à leur pendantes humaines, on peut alors se demander si la puissance de généralisation (et même d'apprentissage) humaine est également discrète. Cette affirmation semble cavalière, car même s'il est possible de supposer que la part consciente de l'esprit humain raisonne à horizon finie et a une puissance dénombrable (transmise d'ailleurs à la machine, apprenant selon des modalités humaines), cette dimension occulte la quantité d'information sans cesse captée et filtrée par notre cerveau ainsi que son assimilation inconsciente, relevant autant de la pensée abstraite que du biologique peut potentiellement générer une puissance de généralisation relevant d'un infini plus large et ainsi fournir une puissance de généralisation continue (relevant davantage de la ligne que du point). Dès lors, comment penser la généralisation chez l'homme alors que, mathématiquement, nos intuitions les plus simples nous font défaut lorsque cette puissance continue intervient (la boule de rayon 1 n'est pas compacte en dimension infinie, RIESZ, 1955)? On peut également se demander si l'extrapolation existe dans de telles structures ou si tout revient à interpoler (HASSON *et al.*, 2020).

Quid de la généralisation en apprentissage machine de nos jours? Qu'espérer alors des réseaux de neurones artificiels et de leur capacité de généralisation relativement à l'humain? Les théorèmes d'approximations universels (voir *e.g.* LU *et al.*, 2017; PARK *et al.*, 2021) assurent que les réseaux de neurones sont capables d'approximer n'importe quelle fonction vivant dans un espace à la puissance du continu (*e.g.* l'espace des fonctions continues à support compact qui n'admet pas de base dénombrable en tant qu'espace de Banach), faisant de ces structures des candidats prometteurs pour appréhender partiellement les mécanismes humains de généralisation. Les approximations prodiguées par ces machines seront, dans un avenir proche, potentiellement suffisamment puissantes pour donner l'illusion d'une capacité de généralisation humaine. Néanmoins, il demeure bon de garder en tête que, si la thèse d'une inégalité fondamentale de nature entre les puissances de généralisation humaine et machine est avérée, alors les réseaux de neurones artificiels n'atteindront jamais pleinement les capacités de compréhension du monde de leurs homologues biologiques. Reste que leur capacité à approximer cette intelligence toute humaine font de ces structures des assistants de valeur, enrichissant les capacités de tout individu. Mieux comprendre la puissance de généralisation machine, être capable de la quantifier, d'identifier les mécanismes qui la

favorisent sont les objets de ce manuscrit.

PREAMBLE: HUMAN LEARNING, MACHINE LEARNING AND GENERALISATION

This manuscript tackles the notion of *generalisation* a notion built upon the general notion of *learning*. For a brief moment, let's take the luxury of forgetting about machines and concentrate on learning at its most human.

Apprehending human learning A human being (here a learner) is structured around experiences, either lived or passed on by others.

The learner then benefits from these experiences in various ways, for instance, by considering a mediated experience to be true (fire burns) and acting according to this. On the contrary, reiteration or denial of this same information may be symptoms of zero truth value. These scenarios can just as easily appear for a lived experience (the question of hallucinations). This first dichotomy in information processing is intrinsically linked to a clearly stated question: does fire burn? Can I trust my senses or have I hallucinated? In these cases, learning has taken place by reducing the intrinsic complexity of an experience to its truth value *w.r.t.* a simple question (in this case with two outcomes). This vision can easily be extended to a finite tree of possibilities through multiple-choice questions. Indeed, we can extend the burning question as follows: what is the intensity of the burn as a function of the temperature of the fire? We can then establish a multitude of answers representing various degrees of burn.

However, many questions cannot be reduced to a finite number of possibilities. For example, what is fire? To answer this question, it is nevertheless possible to exploit multiple facets of experience (wood, twig, rock fire) to propose that fire is the chemical reaction of oxygen in the air with a combustible material, with a supply of energy serving as the trigger.

Then, a legitimate question is: why has mankind understood the nature of fire? This fundamental understanding emerged from practical considerations: how can we stop being cold? Can we eat meat other than raw to reduce the risk of illness? It then takes multiple interactions with the environment to generate experiences and then learn from them to gradually respond to a complex need (how to make a fire to keep yourself warm?).

Thus, through this preliminary analysis, we have found several premises of understanding human learning.

- How is learning formalised structurally? The learner must base the experience on simple questions to acquire primary certainties. These latter acquired, it is possible to reach complex questions by interweaving more and more elementary considerations.
- Where does the need to learn come from? From a practical point of view, the emergence of these complex questions often arises from a relationship between the being and its environment, making it possible to develop contextual objectives. The learner then gradually becomes capable of responding to complex needs through a succession of simple actions.

From human to machine learning Machine learning has been structured around two approaches, the first is symbolic and takes advantage of human extrapolations to teach the machine to manipulate an axiomatic, while the second is statistical, and consists of providing the machine with a large number of experiments so that it learns from multiple empirical examples. We are going to focus on the second approach because it underpins a large part of modern research. This method requires a large number of experiments to be transmitted to the machine, which then extracts the knowledge through optimising procedures. More precisely, the knowledge extracted depends on the question posed and its mathematical translation. We can see parallels with human learning described above: you need experiments and a question to reduce reality to something learnable. To go a step further, the variety of human learning scenarios described above can be applied to modern machine learning: the question "Does fire burn?" can be associated with supervised learning, which learns from multiple-choice questions. The question "What is fire?" can be associated with unsupervised learning, which, in the case of clustering, looks for similarities between numerous experiments that are not induced by the question. Finally, the question "Can I make a fire?" can be linked to reinforcement learning which focuses on the evolution of an agent learning from its interaction with the environment.

From learning to generalisation. Generalisation can be seen as the ability to exploit learning from experience beyond that experience. This encompasses a theoretical and axiomatic understanding of a phenomenon, *i.e.* a fruitful extrapolation, or the ability to exploit the knowledge acquired for a new, yet showing similarities, situations *i.e.* to interpolate experiences.

This dual aspect of generalisation can be found in both humans and machines in a variety of ways. Deep neural networks, which are the spearhead of modern machine learning, are based on finite-dimensional learning spaces, which means that a problem can be learned through a finite number of founding principles. Since the number of principles can be increased as far as numerical capacity allows, we can say that neural networks have discrete generalising power. Given that machine learning methods are

correlated with their human counterparts, we might then ask whether the power of human generalisation (and even learning) is also discrete. This assertion is somewhat bold as even it is assumable that the conscious part of the human mind reasons on a finite horizon and has a discrete generalisation power (transmitted, moreover, to the machine, which learns according to human methods), this dimension obscures the quantity of information constantly captured and filtered by our brain, as well as its unconscious assimilation. In other words, the fact that our brain is as much a part of abstract thought as it is of biological thought can potentially generate a generalisation power that relates to a wider infinity and thus provide a continuous generalisation power (relating more to the line than to the point). So how can we think about generalisation in humans when, mathematically, our simplest intuitions fail us when this continuous power is involved (the ball of radius 1 is not compact in infinite dimension, RIESZ, 1955)? We might also ask whether extrapolation exists in such structures or whether it all boils down to interpolation (HASSON *et al.*, 2020).

What to expect from generalisation in modern machine learning? So what can we expect from artificial neural networks and their ability to generalise to humans? Universal approximation theorems (see *e.g.* LU *et al.*, 2017; PARK *et al.*, 2021) ensure that neural networks are capable of approximating any function living in a space to the power of the continuum (*e.g.* the space of continuous functions with compact support which does not admit a countable base as a Banach space), making these structures promising candidates for partially understanding human generalisation mechanisms. In the near future, machine approximations will potentially be powerful enough to give the illusion of human generalisation capacity. Nevertheless, it is worth bearing in mind that, if the thesis of a fundamental inequality in nature between the powers of human and machine generalisation is confirmed, then artificial neural networks will never fully attain the world-understanding capacities of their human counterparts. It is still worth noticing artificial neural nets ability to approximate this human intelligence makes these structures valuable assistants, enriching the capabilities of any individual. That being said, this manuscript aims to provide a better understanding of generalisation in machine learning, quantifying and indentifying the mechanisms that promote it.

JURY:

Rapporteurs: Alquier (sur)/Chopin (moins)

Membres: Seldin (rapporteur mais peut etre pas fou pour le rapport)/ Pascal Germain (rapporteur) Gérard (si Pierre pas dispo), John Shawe-Taylor (examineur), Emilie Morvant (présidente) + le Maitre + Arnak Dalalyan (trop proche?), Alessandro Rudi
CHALLENGE HERE: being very rigorous on the lit review.

PAC-BAYES LEARNING, A FIELD OF MANY PARADIGMS

1.1 A brief introduction to statistical learning

Statistical learning (VAPNIK, 1999; JAMES *et al.*, 2013) quantifies and identifies how learning algorithms, trained on a specific task using a finite training dataset, generalise to novel, unseen datum. More precisely, a learning agent has to learn how to answer a question, formalised as a *learning problem* being a tuple $(\mathcal{H}, \mathcal{Z}, \ell)$ composed of a *predictor space* on which evolves the agent during the learning process, a *data space* \mathcal{Z} and a *loss function* being the mathematical formulation of the question. Such a minimalistic structure is convenient to encompass a broad range of real-life learning scenarii. To learn, the agent has access to a *training dataset* $\mathcal{S}_m = (\mathbf{z}_i)_{i=1\dots m}$. The most classical way to learn from \mathcal{S}_m is the empirical risk minimisation (ERM), minimising the *empirical risk* $\hat{R}_{\mathcal{S}_m} := \frac{1}{m} \sum_{i=1}^m \ell(h, \mathbf{z}_i)$. In this setting, when \mathcal{S}_m is *i.i.d.* (following the distribution \mathcal{D}), two facets of generalisation are commonly studied in statistical learning for a given agent $h \in \mathcal{H}$.

- First, the *population risk* $R_{\mathcal{D}}(h) := \mathbb{E}_{\mathbf{z} \sim \mathcal{D}}[\ell(h, \mathbf{z})]$ focus on the average performance of our learning agent *w.r.t.* any new situation $\mathbf{z} \in \mathcal{D}$, independent of \mathcal{S}_m , possibly faced by the agent. A small population risk ensure then efficient generalisation.
- Second, the *generalisation gap* $\Delta_{\mathcal{S}_m}(h) := R_{\mathcal{D}}(h) - \hat{R}_{\mathcal{S}_m}(h)$ evaluate the coherence between the empirical risk and the population one. Having a small generalisation gap ensure that the generalisation ability of the agent has the same magnitude than its training performance.

Note that the population risk is a stronger notion of generalisation than the generalisation gap. However, a small generalisation gap (in absolute value) as well as a small empirical risk is enough to ensure a good population risk. Given that modern optimisation algorithm are often enough to ensure a small empirical risk, the generalisation gap has received a particular attention in statistical learning.

Generalisation bounds. Generalisation bounds are inequalities often controlling the generalisation gap by various quantities depending either on \mathcal{H} , \mathcal{Z} or \mathcal{S}_m . We propose below general patterns usually involved in generalisation bounds for an agent $h_{\mathcal{S}_m} \in \mathcal{H}$ depending on \mathcal{S}_m (for instance the output of the ERM).

Expected generalisation bound. For any training set \mathcal{S}_m :

$$\mathbb{E}_{\mathcal{S}_m} [\Delta_{\mathcal{S}_m}(h_{\mathcal{S}_m})] \leq f\left(\text{COMPLEXITY}, \frac{1}{m}\right). \quad (1.1)$$

High-probability generalisation bounds. For any training set \mathcal{S}_m , with probability $1 - \delta$ pver the draw of \mathcal{S}_m :

$$\Delta_{\mathcal{S}_m}(h_{\mathcal{S}_m}) \leq f\left(\text{COMPLEXITY}, \frac{1}{m}, \log \frac{1}{\delta}\right). \quad (1.2)$$

The nature of f and the COMPLEXITY term depend on the facet of the complexity of the learning problem we aim to focus. Celebrated examples are for instance the dimension of \mathcal{H} , if euclidean, the VC dimension of \mathcal{H} (VAPNIK, 2000), the Rademacher complexity (BARTLETT and MENDELSON, 2001, 2002), the stability parameter of a learning algorithm (BOUSQUET and ELISSEEFF, 2000) or the subgaussian diameter of \mathcal{Z} (KONTOROVICH, 2014). Another approach relies on the Bayesian learning paradigm, deriving *posterior* knowledge from data and prior modelling of the environment. Then, the COMPLEXITY can be borrowed from information theory (COVER and THOMAS, 2001), e.g. mutual information (NEAL, 2012), or from optimal transport, e.g. Wasserstein distances (WANG *et al.*, 2019; RODRIGUEZ-GALVEZ *et al.*, 2021).

Those two approaches have various benefits. A notable strength of expected bounds is that they may reach fast convergence rates (*i.e.* faster than $\frac{1}{\sqrt{m}}$) contrary to high-probability one, even when \mathcal{H} is a singleton thanks to the central limit theorem (GRUNWALD *et al.*, 2021). However, expected bounds often involves a theoretical COMPLEXITY which cannot be estimated in practice and may be hard to interpret while high probability bounds may be fully empirical and can be considered with small confidence parameter δ as it is attenuated by a logarithm.

How to choose the complexity term ? An introductory example. There is no evidence proving that a certain notion of complexity is preferable to another. The choice of COMPLEXITY may however be driven by practical considerations, emerging from the learning problem of interest. To illustrate this point, let us focus on the following example, providing two learning problems which differs only from the predictor space \mathcal{H} and which have very different interactions with the VC dimension.

Example 1.1.1. Consider a supervised learning problem where $\mathcal{Z} = \mathbb{R}^k \times \mathcal{Y}$ with $\mathcal{Y} = \{0, 1\}$, k smaller than m and with loss $\ell(h, (x, y)) = \mathbb{1}\{h(x) \neq y\}$. First, assume that \mathcal{H} is the set of linear classifiers; *i.e.* $\mathcal{H}_1 := \{h_\theta(x) = \text{sgn}(\langle \theta, x \rangle)\}$, where $\text{sgn}(a)$ denotes the sign of a . In this case, using the VC dimension may lead to non-vacuous generalisation bounds (VAPNIK, 2000).

1.2. An information-theoretic exposition of PAC-Bayes learning

However, in modern machine learning, deep neural networks are often considered, let us first define a celebrated class of deep neural networks.

Definition 1.1.1. A multilayer perceptron with depth K and architecture $\{N_1, \dots, N_K\}$, denoted as $h_{\mathbf{w}}(\mathbf{x}) := Wh^K(\dots h^1(\mathbf{x})) + b$, is composed of K layers $h^1(\cdot), \dots, h^K(\cdot)$. $W \in \mathbb{R}^{|\mathcal{Y}| \times N_K}$ and $b \in \mathbb{R}^{N_K}$ are the weight matrix and the bias of the last layer, and the i -th layer h^i , composed of N_i nodes, is defined by $h^i(\mathbf{x}) := \sigma_i(W_i \mathbf{x} + b_i)$, where $W_i \in \mathbb{R}^{N_i \times N_{i-1}}$ and the bias $b_i \in \mathbb{R}^{N_i}$ are its weight matrix and bias respectively; $\sigma_i : \mathbb{R}^{N_i} \rightarrow \mathbb{R}^{N_i}$ is an activation function. The weights $\mathbf{w} = \text{vec}(\{W, W_K, \dots, W_1, b, b_K, \dots, b_1\})$ represent the vectorisation of all parameters of the network.

Now, consider the learning problem with the same \mathcal{Z}, ℓ as above, but with \mathcal{H}_2 being the set of multilayer perceptrons *w.r.t.* a fixed depth K and architecture $\{N_1, \dots, N_K\}$. To be consistent with modern practice, assume also that we are in the *overparametrised setting*, meaning that the space \mathcal{H}_2 has a dimension d far greater than m . In this case, VC dimension fails to explain the good generalisation ability (seen in practice) of multilayer perceptrons (BARTLETT and MAASS, 2003).

Understanding the generalisation ability of deep neural networks remains nowadays a major challenge and in what follows, we focus on a modern branch of learning theory which provided non-vacuous bounds of the generalisation ability of deep neural networks: PAC-Bayes learning.

1.2 An information-theoretic exposition of PAC-Bayes learning

PAC-Bayes learning is a recent branch of learning theory which emerged in the late 90s via the seminal work of (SHAWE-TAYLOR and WILLIAMSON, 1997; MCALLESTER, 1998, 1999, 2003) and later pursued by (CATONI, 2003, 2007). Modern surveys recently emerged to describe the various advances in the field (GUEDJ, 2019; HELLSTRÖM *et al.*, 2023; ALQUIER, 2024). Similarly to the various subfields of statistical learning described in Section 1.1, PAC-Bayes theory is designed to provide generalisation bounds involving a COMPLEXITY term apprehending a facet of the complexity of the learning problem. In PAC-Bayes, this term is inspired from the Bayesian learning paradigm of designing a *posterior* knowledge of the learning problem based on the positive impact of data onto a *prior* knowledge of the considered situation. A concrete example of Bayesian learning would be an explorer mapping an ill-known territory. The explorer has to adapt the existing maps at its disposal before exploration to its discov-

eries, generating a new map imbricating the benefits of both the prior one alongside its findings. From a mathematical perspective, the Bayes approach relies on the Bayes formula, providing an update recipe from a prior distribution $P \in \mathcal{M}(\mathcal{H})$ over the predictor space \mathcal{H} to a posterior $Q \in \mathcal{M}(\mathcal{H})$ through a likelihood. On the contrary, PAC-Bayes, while inspired from the Bayesian philosophy, does not relies on the Bayes formula but instead on tools from information theory. This general approach benefits from additional flexibility as PAC-Bayes can be linked and applied to Bayesian learning (see GUEDJ, 2019) but also blurs the notion of prior and posterior distributions, now independent of the fundamental Bayes formula. We further develop those points through two celebrated bounds: the McAllester and Catoni ones.

Two fundamental results

The McAllester's bound (MCALLESTER, 2003) enriched with Maurer's trick (MAURER, 2004) and Catoni's bound (ALQUIER *et al.*, 2016, Theorem 4.1, being a relaxation of CATONI, 2007, Theorem 1.2.6) are probably the most known high-probability PAC-Bayes bounds. We recall them in Proposition 1.2.1.

Proposition 1.2.1. Assume \mathcal{S}_m to be *i.i.d.*

McAllester's bound, (Maurer, 2004, Theorem 5). For any $\delta \in (0, 1)$, $\ell \in [0, 1]$, any data-free prior $P \in \mathcal{M}(\mathcal{H})$, with probability at least $1 - \delta$, for any posterior $Q \in \mathcal{M}(\mathcal{H})$,

$$\Delta_{\mathcal{S}_m}(Q) \leq \sqrt{\frac{\text{KL}(Q, P) + \ln \frac{2\sqrt{m}}{\delta}}{2m}}. \quad (1.3)$$

Catoni's bound, (Alquier *et al.*, 2016, Theorem 4.1). For any $\lambda \in \mathbb{R}/\{0\}$, $\delta \in (0, 1)$, ℓ being σ^2 -subgaussian and a data-free prior P , with probability at least $1 - \delta$ over \mathcal{S} , for any $Q \in \mathcal{M}(\mathcal{H})$,

$$\Delta_{\mathcal{S}_m}(Q) \leq \frac{\text{KL}(Q, P) + \log(1/\delta)}{\lambda} + \frac{\lambda\sigma^2}{2m}. \quad (1.4)$$

For both results, $\Delta_{\mathcal{S}_m}(Q)$ denotes the expected generalisation gap *w.r.t.* Q and KL denotes the Kullback-Leibler divergence.

Recall that a random variable X is σ^2 -subgaussian if for any $\lambda \in \mathbb{R}$, $\mathbb{E}[\exp(\lambda(X - \mathbb{E}[X]))] \leq \exp\left(\frac{\lambda^2\sigma^2}{2}\right)$ and that any loss $\ell \in [0, C]$ is C -subgaussian. Both McAllester and Catoni bounds fit the general shape of (1.2). In both cases, $\text{COMPLEXITY} = \text{KL}(Q, P)$ and f varies. The immediate link with the Bayesian philosophy of learning

1.2. An information-theoretic exposition of PAC-Bayes learning

is that the prior has to be data-free. However, (1.3) and (1.4) are both valid simultaneously for any posterior, which is strictly more general than considering the Bayesian posterior. Note that if λ is optimised, then Catoni's bound would boil down to an upgraded McAllester bound without the $\log(\sqrt{m})$ term, but such an optimisation is not feasible as λ has to be chosen independently of the dataset \mathcal{S}_m . Note that this gap has been recently filled by DUPUIS and ŞİMŞEKLI (2024, Theorem 33). While the theoretical links between those two bounds are clear, they involve two different toolboxes: McAllester's bound heavily relies on the KL divergence between Bernoullis alongside calculation tricks exploiting the boundedness of the loss while the original Catoni's bound (CATONI, 2007, Theorem 1.2.6) exploits tools from statistical physics. The relaxation (1.4) proposed here is reachable by a few key arguments, involved in a vast majority of PAC-Bayes proofs. We propose it below for pedagogical purpose.

Proof of Equation (1.4). Note that the first part of the proof holds for a large part of PAC-Bayes literature. **A generic pattern for PAC-Bayes bounds.** This part is designed upon two cornerstones, retrievable in many existing results: the change of measure inequality (CSISZÁR, 1975; DONSKER and VARADHAN, 1976 – see also BANERJEE, 2006; GUEDJ, 2019 for a proof) and Markov's inequality.

Lemma 1.2.1 (Change of measure inequality). For any measurable function $\psi : \mathcal{H} \rightarrow \mathbb{R}$ and any distributions Q, P on \mathcal{H} :

$$\mathbb{E}_{h \sim Q}[\psi(h)] \leq \text{KL}(Q, P) + \log(\mathbb{E}_{h \sim P}[\exp(\psi(h))]).$$

For a given $\lambda > 0$, the change of measure inequality is then applied to a certain function $f_m : \mathcal{H} \rightarrow \mathbb{R}$, possibly involving \mathcal{S}_m : for all posteriors Q ,

$$\mathbb{E}_{h \sim Q}[f_m(h)] \leq \text{KL}(Q, P) + \log(\mathbb{E}_{h \sim P}[\exp(f_m(h))]). \quad (1.5)$$

To deal with the random variable $X(\mathcal{S}_m) := \mathbb{E}_{h \sim P}[\exp(f_m(h))]$, our second building block is Markov's inequality ($\mathbb{P}(X > a) \leq \frac{\mathbb{E}[X]}{a}$) which we apply for a fixed $\delta \in (0, 1)$ on $X(\mathcal{S}_m)$ with $a = \mathbb{E}_{\mathcal{S}_m}[X(\mathcal{S}_m)]/\delta$. Taking the complementary event gives that for any m , with probability at least $1 - \delta$ over the sample \mathcal{S}_m , $X(\mathcal{S}_m) \leq \mathbb{E}_{\mathcal{S}_m}[X(\mathcal{S}_m)]/\delta$, thus:

$$\mathbb{E}_{h \sim Q}[f_m(h)] \leq \text{KL}(Q, P) + \log(1/\delta) + \log(\mathbb{E}_{h \sim P} \mathbb{E}_{\mathcal{S}_m}[\exp(f_m(h))]). \quad (1.6)$$

Note that in (1.6), we swapped the two expectations in the last term thanks to Fubini's theorem and the fact that P is data-free.

Proving Catoni's bound. Now, we take $f_m(h) = \lambda \Delta_{\mathcal{S}_m}$ and consider for any $h \in \mathcal{H}$, $A(h) = \mathbb{E}_{\mathcal{S}_m}[\exp(f_m(h))]$.

Note that, given \mathcal{S}_m is iid,

$$A(h) = \prod_{i=1}^m \mathbb{E}_{\mathcal{S}_m} \left[\exp \left(\frac{\lambda}{m} (R_{\mathcal{D}}(h) - \ell(h, \mathbf{z}_i)) \right) \right],$$

and thanks to Hoeffding's lemma alongside ℓ being σ^2 -subgaussian,

$$A(h) \leq \prod_{i=1}^m \exp \left(\frac{\lambda^2 \sigma^2}{2m^2} \right) = \exp \left(\frac{\lambda^2 \sigma^2}{2m} \right).$$

Plugging this upper bound in (1.6) and dividing by λ concludes the proof. ■

The generic pattern (1.6), allows to retrieve many PAC-Bayes bounds, starting with McAllester's one, where $f_m = kl(R_{\mathcal{D}}(h), \hat{R}_{\mathcal{S}_m}(h))$, kl being the KL divergence between Bernoullis and completing with the subtle calculations of MAURER (2004). This pattern is also valid, for instance, for the results of GERMAIN *et al.* (2009), the Bernstein PAC-Bayesian bounds of TOLSTIKHIN and SELDIN (2013) and MHAMMEDI *et al.* (2019) and many other results, *e.g.* THIEMANN *et al.* (2017), GUEDJ and ROBBIANO (2018), HOLLAND (2019), and WU and SELDIN (2022). This then pins two major points for a large part of PAC-Bayes literature:

1. Interpreting PAC-Bayes from a Bayesian prism is legitimated by the change of measure inequality, yet the KL divergence. More generally, this property allows interpreting PAC-Bayes under a more general information theoretic paradigm, where information from the prior is partially transferred to the posterior (here by absolute continuity to keep the KL finite). This vision encompasses the Bayesian one, while being less restrictive.
2. The statistical properties of the learning problem are linked to the exponential moment coming from the change of measure inequality, this often implies the strong assumptions of Proposition 1.2.1: data-free prior, bounded or subgaussian losses (sometimes attenuated to subexponentiality CATONI, 2004).

A theory suited for Example 1.1.1? The two previous points show that Proposition 1.2.1 holds for learning problem with light-tailed losses (often bounded), *i.i.d.* data, encompassing classification tasks for instance. Then, PAC-Bayes learning seems suited to understand, on such problems, the McAllester and Catoni bounds are suited to the learning problem $(\mathcal{H}_2, \mathcal{Z}, \ell)$ of Example 1.1.1.

However, the question of their tightness is unsolved as we do not know the behavior of the KL term in practice. Furthermore the question of which distribution Q should be taken in Proposition 1.2.1 remains open.

1.3 From theory to learning algorithms

Algorithms associated to McAllester and Catoni bounds

A shared particularity of McAllester and Catoni bounds is that they are both fully empirical. Then it is possible to minimise them in practice and thus, deriving new theory-driven learning algorithms which are expected to have at worse, a small generalisation gap and at best, a small population risk. More precisely, learning algorithms associated to Proposition 1.2.1 are stated below:

$$Q_M := \operatorname{argmin}_{Q \in \mathcal{C}} \hat{R}_{S_m}(Q) + \sqrt{\frac{\text{KL}(Q, P)}{2m}}. \quad (1.7)$$

For any $\lambda > 0$,

$$Q_C := \operatorname{argmin}_{Q \in \mathcal{C}} \hat{R}_{S_m}(Q) + \frac{\text{KL}(Q, P)}{\lambda}. \quad (1.8)$$

In both cases, $\mathcal{C} \subseteq \mathcal{M}(\mathcal{H})$ is the class of distributions on which we optimise. The choice of \mathcal{C} can come from a priori knowledge of the problem or from optimisation concerns to make the KL divergence tractable.

Knowing Catoni's bound is a relaxation of McAllester's one, it seems more natural to consider Q_M over Q_C . However, the presence of a square root in (1.7) can be challenging for practical optimisation. We illustrate this below.

Example 1.3.1. Consider the case where, for a given $\sigma > 0$, $\mathcal{C} = \{\mathcal{N}(\mu, \sigma^2 \text{Id}) \mid \mu \in \mathbb{R}^d\}$. Then for any $P = \mathcal{N}(\mu_1, \sigma^2 \text{Id})$, $Q = \mathcal{N}(\mu_2, \sigma^2 \text{Id})$, $\text{KL}(Q, P) = \frac{\|\mu_1 - \mu_2\|^2}{2\sigma^2}$. Then, optimising (1.7) in this case implies to lose the strong convexity of the KL divergence while it is retained for (1.8).

Another practical advantage of (1.8) over (1.7) emerges when $\mathcal{C} = \mathcal{M}(\mathcal{H})$. In this case, Catoni's bound admits a closed form solution, while McAllester's one should be numerically optimised on all the space of distributions, which is not feasible. This closed form, extracted from CATONI (2003, Section 5.1), is recalled below.

$$\text{When } \mathcal{C} = \mathcal{M}(\mathcal{H}), \quad dQ_C(h) = \frac{\exp(-\lambda \hat{R}_{S_m}(h))}{\mathbb{E}_{h \sim P}[\exp(-\lambda \hat{R}_{S_m}(h))]} dP(h) \quad (1.9)$$

Then, $Q_C = P_{-\lambda \hat{R}_{S_m}}$ is the *Gibbs posterior* associated to $P, \lambda \hat{R}_{S_m}$. By introducing Gibbs posterior in statistical learning, CATONI (2007) draws a theoretical link between statistical physics and learning theory. Unfortunately, Gibbs posteriors often requires Monte Carlo methods to be implemented, which can be in practice time-consuming. Below we then focus on PAC-Bayes algorithms working on a subset \mathcal{C} of $\mathcal{M}(\mathcal{H})$.

Instantiation and efficiency of PAC-Bayesian algorithms

A general pattern for PAC-Bayesian algorithms The introductory examples (1.7),(1.8) unveil a general design for any KL-based PAC-Bayesian algorithm. satisfy a tradeoff between (i) the empirical risk, showing that the learner has to fit the training dataset, and (ii) a *regulariser* being a function of $\text{KL}(Q, P)$. This regulariser ensures that, during training, the learner will not overfit on training data. This training ensures a good generalisation ability as long as the associated generalisation bound is small.

We then understand better the conceptual ins and outs of PAC-Bayes algorithms, two unanswered questions remains:

1. How are those algorithms instantiated in practice?
2. Are these algorithms efficient and do they come with non-vacuous theoretical guarantees?

Instantiating a PAC-Bayes algorithm In practice, using a single prior P usually does not work but it remains theoretically possible to consider a finite set of priors. Indeed, if one wants to consider k priors, then it is possible to consider k PAC-Bayes bounds holding for each of those priors with probability at least $1 - \frac{\delta}{k}$ and then consider a union bound, such a set of priors is called a *grid*. This method has been widely used in many PAC-Bayes work with clever grids, deteriorating initial bounds at the cost of supplementary $\log(n)$ or $\log \log(n)$ (divided by m), see *e.g.* ALQUIER (2024). Similarly, for Catoni-typed algorithms where a parameter λ is involved, it is also possible to consider grids for it. In both cases, considering grids allows to optimise on both the prior, the posterior (and possibly λ when involved) and then taking the closest value of those optimised parameters on the grid to still obtain theoretical guarantees. Another technique to ensure a good prior is to sacrifice a part of the training set to pre-train P . Doing so, the prior is then data-dependent and yields tighter bounds alongside increased performance (PEREZ-ORTIZ *et al.*, 2021c).

Efficiency of PAC-Bayes algorithms on supervised learning problems. The work of DZIUGAITE and ROY (2017) showed that optimising (1.7) when \mathcal{C} is a class

1.3. From theory to learning algorithms

of Gaussian measures for the weights of a deep neural network yields non-vacuous generalisation bound, meaning that the generalisation benefits of PAC-Bayesian training on deep nets can be theoretically ensured. In their work, authors used the instantiation toolbox described above, alongside a preliminary use of Stochastic Gradient Descent (SGD) to update Q before the PAC-Bayes training algorithm. This promising work paved the way to various extensions, providing non-vacuous guarantees for a wide range PAC-Bayes algorithms (LETARTE *et al.*, 2019; RIVASPLATA *et al.*, 2019; DZIUGAITE *et al.*, 2021; PEREZ-ORTIZ *et al.*, 2021a,b,c; BIGGS and GUEDJ, 2022a, 2023), showing that the PAC-Bayes toolbox provides elements of answer to understand the generalisation ability of neural networks. Beyond generalisation guarantees, PAC-Bayes bounds may also be useful to propose original training methods, even if the associated guarantees are vacuous (BIGGS and GUEDJ, 2021, 2022b). Another important empirical use is to exploit PAC-Bayes bounds as correlation measures, to see whether a decrease of the bound is related to an increased generalisation ability of the learner. For instance NEYSHABUR *et al.* (2017) used McAllester's bound (1.3) as a 'flatness' measure and showed that it correlates well with a good generalisation ability for a few learning problems. This conclusion has been extended to a wider range of problems in DZIUGAITE *et al.* (2020) and JIANG *et al.* (2020).

PAC-Bayes algorithms beyond supervised learning. While supervised learning is a widely used setting to perform experiments in PAC-Bayes learning (often involving celebrated datasets such as MNIST or CIFAR-10), the McAllester bound holds for any learning problem with bounded loss, going beyond this setting. This theoretical flexibility has been exploited to derive PAC-Bayesian algorithm for various learning settings reinforcement learning (FARD and PINEAU, 2010), multi-armed bandits (SELDIN *et al.*, 2011, 2012b; SAKHI *et al.*, 2023), meta-learning (AMIT and MEIR, 2018; DING *et al.*, 2021; FARID and MAJUMDAR, 2021; ROTHFUSS *et al.*, 2021, 2022) to name but a few.

Is Example 1.1.1 tackled now? (DZIUGAITE and ROY, 2017) and following works have provided a positive answer by obtaining non-vacuous guarantees (sometimes tight) for $(\mathcal{H}_2, \mathcal{Z}, \ell)$ of Example 1.1.1 for various \mathcal{Z} (being, e.g., set of images for MNIST CIFAR-10 etc...). To obtain such guarantees, a PAC-Bayesian training needs to be performed to minimise its associated theoretical bound. That being said, several questions then legitimately emerge.

- Modern machine learning often implies learning problems where assumptions such as bounded (or subgaussian) losses or *i.i.d.* data do not hold. Is PAC-Bayes theory extendable beyond those assumptions?

- As shown in DZIUGAITE and ROY (2017), the PAC-Bayesian training is often combined to another procedure (e.g. ERM) to yield non-vacuous bounds. However, PAC-Bayes bounds do not bring the theoretical understanding of such additional methods, often outputting deterministic predictors (*i.e.* Dirac distributions). This kind of predictor is not allowed in (1.3), (1.4). Is it possible to obtain PAC-Bayes bounds valid for such methods?

1.4 An optimisation perspective of PAC-Bayes

The questions raised at the end of the previous part are important as they underline a gap between the information-theoretic approach of PAC-Bayes bounds and practical optimisation. A supplementary example of this is the grid required in practice to optimise the prior (and/or λ in Catoni's bound). Indeed this hybrid solution is required to roughly fit theory, (exploiting a single prior) and practice (optimising freely the prior on a continuous space), while not being truly adapted to any of these settings. This then raises the following fundamental question:

Can we think PAC-Bayes learning from an optimisation perspective?

The elements of answer to this question are multiple. First, one can mix PAC-Bayes argument with geometric properties of optimisation procedure to obtain generalisation bounds designed for specific algorithms exploiting their geometric properties and assumptions, including but not limited to, SGD, Langevin dynamics (LONDON, 2017; DZIUGAITE and ROY, 2018a; NEU *et al.*, 2021; CLERICO *et al.*, 2022; HAGHIFAM *et al.*, 2023; ZHOU *et al.*, 2023). Those works shows both convergence properties as well as minimax rates, showing the impact of PAC-Bayes learning to provide a better theoretical understanding of the generalisation ability of concrete algorithms.

A second approach consists in describe general principles that should be satisfied by the various terms and assumptions in PAC-Bayes when looking at this through the prism of optimisation. We propose such an analysis below.

An optimisation-driven view of PAC-Bayes

- **Statistical assumptions.** While ℓ satisfies desirable geometric properties (convexity, gradient lipschitz ...), no statistical assumption is needed. However, the output of two runs of a stochastic optimisation algorithm on the same training set may vary a lot. For the specific case of SGD, many works shows that SGD exhibits a heavy-tailed behavior (see e.g. ŞİMŞEKLI *et al.*, 2019; ZHANG *et al.*, 2020; GÜRBÜZBALABAN *et al.*, 2021). Given that such heavy-tailed consideration emerges from the most celebrated algorithm, generalisation bounds, from an optimisation point of view, should hold with weak statistical assumptions on the dataset, even if additional geometric assumptions on the loss may emerge.

- **The role of the prior.** The information-theoretic approach justify the Bayesian view of the prior as discussed earlier. In this spirit it is also possible to sacrifice a part of the training set (*i.e.* of the available information) to enrich P and then partially explain the efficiency of an information-theoretic training. Those two vision are not easily linked to optimisation concerns as the first one would be linked to a 'good' initialisation, something we cannot know in advance, while the second makes little sense as P is obtained through a first, unexplained, optimisation process which is necessary to understand the efficiency of the second part of training, outputting Q . From an optimisation stance, we suggest to assign only two possible roles to P : (i) the initialisation of the optimisation algorithm, then its impact should be attenuated through the learning process and (ii) a minimiser we aim to reach through optimisation. In this case, its impact is crucial as it translates the speed of convergence of our learning algorithms.
- **The place of stochastic predictors.** Involving a KL divergence as a complexity brings a particular focus on stochastic predictors, drawn from a distribution Q . Classical PAC-Bayes bounds usually focus on the average performance of such a predictor (hence the expectation over Q in (1.3),(1.4)), but recent extensions directly proposed guarantees for a single draw over Q (RIVASPLATA *et al.*, 2020; VIALARD *et al.*, 2023a). However, involving a KL implies that Q has to be absolutely continuous *w.r.t.* P , meaning that the support of Q cannot go beyond the one of P : this excludes the case of Dirac distributions, *i.e.* deterministic predictors. This is a clear limitation of the information-theoretic approach, as many learning algorithms outputs a deterministic predictor and thus, should be avoided to be in line with common practice in optimisation.

Those three points, while not necessarily considered explicitly through the lens of optimisation have been recently challenged, and modern advances are gathered below.

PAC-Bayes beyond the usual setting

Recall that according to what we saw in McAllester's bound (1.3) and Catoni's one (1.4), we denote by usual setting a bound holding for *i.i.d.* \mathcal{S}_m , with bounded or subgaussian losses and involving a KL divergence as COMPLEXITY term. Many works overcame at least one of this assumption as precised below.

Beyond *i.i.d.* data The work of FARD and PINEAU (2010) established links between reinforcement learning and PAC-Bayes theory. This naturally led to the study of PAC-Bayesian bound for martingales instead of *i.i.d.* data (SELDIN *et al.*, 2011, 2012a,b). Also, PAC-Bayesian bound for lifelong learning (PENTINA and LAMPERT, 2014; FLYNN *et al.*, 2022) challenged also the *i.i.d.* assumption. We also denote that

the PAC-Bayes bound for meta learning (AMIT and MEIR, 2018; DING *et al.*, 2021; FARID and MAJUMDAR, 2021; ROTHFUSS *et al.*, 2021, 2022) consider independent but non-identically distributed datasets (corresponding to different tasks).

Avoiding light-tailed losses. Light-tailed losses encompass bounded, subgaussian, subexponential losses. Deriving PAC-Bayes bound for heavy-tailed losses, starting from AUDIBERT and CATONI (2011) which provided PAC-Bayes bounds for least square estimators with heavy-tailed random variables. Their results was suboptimal with respect to the intrinsic dimension and was followed by further works from CATONI (2016) and CATONI and GIULINI (2017). More recently, this question has been addressed in the works of ALQUIER and GUEDJ (2018), HOLLAND (2019), KUZBORSKIJ and SZEPESVÁRI (2019), and HADDOUCHE *et al.* (2021), extending PAC-Bayes to heavy-tailed losses under additional technical assumptions.

Towards data-dependent priors. The work of (CATONI, 2007; LEVER *et al.*, 2010, 2013) proposed priors, not directly data-dependent, but depending of the data distribution \mathcal{D} when *i.i.d.* data are considered. can be informed by the data-generating distribution, PARRADO-HERNÁNDEZ *et al.* (2012), ONETO *et al.* (2016), DZIUGAITE and ROY (2017), and MHAMMEDI *et al.* (2019) also obtained PAC-Bayes bound with data-dependent priors by infusing directly data in the prior (and sacrificing a part of the dataset). The drawback of this method is that, in practice, such a prior allows tighter bounds, but at the cost of a reduced theoretical understanding as the prior is in practice often learned via ERM, and the PAC-Bayes bound hardly gives insights on what happens during this pre-training. Furthermore, if this pre-training has already made the bound converge to a minimum generalising well, then the PAC-Bayes training has no effect and the associated bound is no more than a test bound (the case $Q = P$). It has been shown for instance in (PEREZ-ORTIZ *et al.*, 2021a) that when P is trained with a consequent fraction of data, then the generalisation performance of the pre-trained P was roughly the same than Q , obtained from P after a PAC-Bayesian training. It is then unclear how impacting are PAC-Bayes methods compared to a test bound in this case. To alleviate this issue, another original route (DZIUGAITE and ROY, 2018b) exploits differential privacy to replace the data-free prior by a differentially private one, making possible to consider the prior as the learning objective (in their case a Gibbs posterior).

Beyond KL divergence. Several works allowed to extend PAC-Bayes beyond KL divergences. The most investigated route is to focus on the more general class of f -divergence, which include, e.g., KL, χ^2 , Rényi divergences among others (ALQUIER and GUEDJ, 2018; OHNISHI and HONORIO, 2021; PICARD-WEIBEL and GUEDJ, 2022; VIALARD *et al.*, 2023a). However, f -divergences still implies absolute conti-

nuity of Q *w.r.t.* P . Another route recently emerged (AMIT *et al.*, 2022), replacing f -divergences by integral probability metrics (IPMs), finally allowing Dirac distribution in PAC-Bayes.

These works have sometimes been explicitly driven by optimisation considerations (DZUGAITE and ROY, 2018b involved differential privacy to numerically tighten their bound without sacrificing data). However, in many cases, the information-theoretic vision of PAC-Bayes remained majoritary. In what follows, the contributions of this manuscript are designed *w.r.t.* the optimisation view of PAC-Bayes detailed above.

1.4.1 Contributions of this thesis

Detail broadly what generalisation is, to what kind of structures it is applied (neural nets or linear classifier eg). Details on the other hand what optimisation is doing (ERM eg) and explain that interestingly in various methods, reaching minimisers of empirical objectives is enough to ensure a good generalisation ability. From this, discuss about the current limitations of generalisation: either not going so often beyond light-tailed assumptions or noticing that the interplays between generalisation (statistical arguments) and optimisation (geometric ones) remains uncharted for a vast range of cases.

Vision: after generic paragraphs on generalisation and optimisation, do a broader paragraph on PAC-Bayes and details the problem of existing PAC-Bayes approach:

Says that PAC-Bayes spontaneously offer a clear link from generalisation to optimisation by providing new learning algorithms: this implicitly suggests assumptions on the loss (eg convex) or on the regulariser (KL between Gaussians to get a strongly convex function) to make sure the minimisation goes well and thus build a bridge with optimisation.

TODO look if there are links from optimisation to PAC-Bayes (must have been some with Dziugaite, Neu with SGD).

Here, we are studying the interplays on both directions. First, we take the opposite perspective and, starting from optimisation benefits/perspectives, we want to understand generalisation, to do so we have several routes within PAC-Bayes. Second, we investigate deeper on the influence of generalisation bounds to derive novel learning algorithms

Thinking the role of the prior in PAC-Bayes: in a similar manner than initialisation/goal to attain in optimisation: if we target a data-free posterior (eg Gibbs catoni) then ok: we target the learning objective. Otherwise, it is common to compare to a random initialisation point of a learning procedure: meaningless. Answer: Online PAC-Bayes which allows, among other, to make the prior evolve through time. (note that there is also either the data-dependent prior: its bad and the differential privacy approach, which is nice)

Switching from statistical to geometric assumptions on the loss. Most of the bounds holds for data-free light-tailed losses: do not necessarily fit the reality of losses involved in optimisation, often unbounded, and either convex, gradient Lipschitz or smooth. Answer: PAC-Bayes for heavy-tailed martingales and flat minima.

Can the convergence properties of optimisation procedure play a role in generalisation? Direct answer: Wasserstein PAC-Bayes, Indirect one: flat minima.

Can we derive generalisation-based learning algorithms beyond Gaussian or Gibbs distributions? Answer: Yes as a byproduct: both in Online PAC-Bayes, Flat Minima and Paper with Paul

Precise the structure of the document: see it as a natural flow where one question implies another one:

1.4. An optimisation perspective of PAC-Bayes

Light-tailed losses? \rightarrow supermartingales!

But then : what should we do of the prior? \rightarrow if you see it as an initialisation: Online PAC-Bayes with novel online algorithms or Flat Minima to attenuate the impact of the prior through fast rates.

What if it should be the optimisation goal? \rightarrow Wasserstein PAC-Bayes

CHALLENGE HERE: being very rigorous on the lit review.

PAC-BAYESIAN BOUNDS FOR MARTINGALES AND HEAVY-TAILED LOSSES

2

This chapter is based on the following paper

MAXIME HADDOUCHE and BENJAMIN GUEDJ. PAC-Bayes Generalisation Bounds for Heavy-Tailed Losses through Supermartingales. *Transactions on Machine Learning Research*. (2023)

Abstract

TODO: put general bounds for martingales and batch learning as corollary

MITIGATING INITIALISATION IMPACT BY REAL-TIME CONTROL: ONLINE PAC-BAYES LEARNING

This chapter is based on the following papers

MAXIME HADDOUCHE and BENJAMIN GUEDJ. Online PAC-Bayes Learning. *Advances in Neural Information Processing Systems (NeurIPS)*. (2022)

MAXIME HADDOUCHE and BENJAMIN GUEDJ. PAC-Bayes Generalisation Bounds for Heavy-Tailed Losses through Supermartingales. *Transactions on Machine Learning Research*. (2023)

PAUL VIALARD, MAXIME HADDOUCHE, UMUT SIMSEKLI, and BENJAMIN GUEDJ. Learning via Wasserstein-Based High Probability Generalisation Bounds. *Advances in Neural Information Processing Systems (NeurIPS)*. (2023b)

Contents

4.1	Introduction	42
-----	------------------------	----

Abstract

Put OPB here. Precise in the intro that the martingale bounds allow to go beyond batch learning but that this has never been made for OL. Put the supermartingale OPB bound in a supplementary section and the Online WPB bound after the main results of OPB to reach heavy-tailed losses.

MITIGATING INITIALISATION IMPACT THROUGH FLAT MINIMA: FAST RATES FOR SMALL GRADIENTS

This chapter is based on the following paper

TODO

Contents

5.1	Introduction	44
-----	------------------------	----

Abstract

This is the PLS paper, precise that the supermartingales bounds are richer than simply recovering classical batch guarantees: we can incorporate gradient norms, which explains generalisation when a flat minima is reached.

4.1 Introduction

WASSERSTEIN PAC-BAYES LEARNING: EXPLOITING OPTIMISATION GUARANTEES TO EXPLAIN GENERALISATION

This chapter is based on the following papers

MAXIME HADDOUCHE and BENJAMIN GUEDJ. Wasserstein PAC-Bayes Learning: A Bridge Between Generalisation and Optimisation. *arXiv*. abs/2304.07048. (2023)
PAUL VIALARD, MAXIME HADDOUCHE, Umut SIMSEKLI, and BENJAMIN GUEDJ. Learning via Wasserstein-Based High Probability Generalisation Bounds. *Advances in Neural Information Processing Systems (NeurIPS)*. (2023b)

Contents

A.1	JENSEN's Inequality	45
A.2	MARKOV's Inequality	45
A.3	Ville's Inequality	46

Abstract

Put WPB here, precise that, when the prior is seen as the learning goal, it is possible for a certain optimisation algorithm to directly incorporate sound geometric optimisation guarantee into a generalisation bound, trading the hope to reach a flat minima with a sound convergence guarantees. However, this comes at the cost of the explicit impact of the dimension. Also put the paper with Paul(batch bounds) as a supplementary content.

5.1 Introduction

SOME MATHEMATICAL TOOLS

A.1 Jensen's Inequality

Theorem A.1.1 (JENSEN's Inequality). Let X a random variable following a probability distribution ν with f a real-valued measurable convex function, we have

$$f\left(\mathbb{E}_{X \sim \nu}[X]\right) \leq \mathbb{E}_{X \sim \nu}[f(X)].$$

Proof. Since $f()$ is a convex function, the following inequality holds, i.e., we have

$$\forall X', \quad a\left(X' - \mathbb{E}_{X \sim \nu}[X]\right) \leq f(X') - f\left(\mathbb{E}_{X \sim \nu}[X]\right),$$

where a is the tangent's slope. By taking the expectation to both sides of the inequality, we have

$$\underbrace{a\left(\mathbb{E}_{X \sim \nu}[X] - \mathbb{E}_{X \sim \nu}[X]\right)}_{=0} \leq \mathbb{E}_{X \sim \nu}[f(X)] - f\left(\mathbb{E}_{X \sim \nu}[X]\right).$$

Hence, by rearranging the terms, we prove the claimed result. ■

A.2 Markov's Inequality

Theorem A.2.1 (MARKOV's Inequality). Let X a non-negative random variable following a probability distribution ν and $\delta > 0$, we have

$$\mathbb{P}_{X \sim \nu}[X \geq \delta] \leq \frac{\mathbb{E}_{X \sim \nu}[X]}{\delta}.$$

Proof. First of all, remark that we have the following inequality for any X

$$\delta \mathbb{1}[X \geq \delta] \leq X \mathbb{1}[X \geq \delta] \leq X. \quad (\text{A.1})$$

Indeed, on the one hand, if $X < \delta$, $\mathbb{1}[X \geq \delta] = 0$, the inequality holds trivially. On the other hand, if $X \geq \delta$, $\mathbb{1}[X \geq \delta] = 1$ and the inequality becomes $\delta \leq X$, which is true. By taking the expectation of Equation (A.1), we have

$$\mathbb{E}_{X \sim \nu} [\delta \mathbb{1}[X \geq \delta]] \leq \mathbb{E}_{X \sim \nu} [X].$$

From the fact that the expectation of a constant is the constant and by definition of the probability, we have

$$\delta \mathbb{P}_{X \sim \nu} [X \geq \delta] \leq \mathbb{E}_{X \sim \nu} [X] \iff \mathbb{P}_{X \sim \nu} [X \geq \delta] \leq \frac{\mathbb{E}_{X \sim \nu} [X]}{\delta},$$

which is the desired result. ■

A.3 Ville's Inequality

Lemma A.3.1 (Ville's maximal inequality for supermartingales). Let (\mathcal{F}_t) be a filtration and (Z_t) a non-negative super-martingale satisfying $Z_0 = 1$ a.s. If Z_t is adapted to \mathcal{F}_t and $\mathbb{E}[Z_t | \mathcal{F}_{t-1}] \leq Z_{t-1}$ a.s., $t \geq 1$, then, for any $0 < \delta < 1$, it holds

$$\mathbb{P}(\exists T \geq 1 : Z_T > \delta^{-1}) \leq \delta.$$

Proof. We apply the optional stopping theorem (DURRETT, 2019, Thm 4.8.4) with Markov's inequality defining the stopping time $i = \inf\{t \geq 1 : Z_t > \delta^{-1}\}$ so that

$$\mathbb{P}(\exists t \geq 1 : Z_t > \delta^{-1}) = \mathbb{P}(Z_i > \delta^{-1}) \leq \mathbb{E}[Z_i] \delta \leq \mathbb{E}[Z_0] \delta \leq \delta.$$
■

ADDITIONAL BACKGROUND

B

B.1 Online Learning

B.2 Wasserstein distances and optimisation on distributions spaces

B.3 Differential privacy

B.4 Log-Sobolev and Poincaré Inequalities

REFERENCES

PIERRE ALQUIER. User-friendly Introduction to PAC-Bayes Bounds. *Foundations and Trends® in Machine Learning*. (2024)

—— Cited on pages 23, 28.

PIERRE ALQUIER and BENJAMIN GUEDJ. Simpler PAC-Bayesian bounds for hostile data. *Machine Learning*. (2018)

—— Cited on page 32.

PIERRE ALQUIER, JAMES RIDGWAY, and NICOLAS CHOPIN. On the properties of variational approximations of Gibbs posteriors. *Journal of Machine Learning Research*. (2016)

—— Cited on page 24.

RON AMIT, BARUCH EPSTEIN, SHAY MORAN, and RON MEIR. Integral Probability Metrics PAC-Bayes Bounds. *Advances in Neural Information Processing Systems (NeurIPS)*. (2022)

—— Cited on page 33.

RON AMIT and RON MEIR. Meta-Learning by Adjusting Priors Based on Extended PAC-Bayes Theory. *International Conference on Machine Learning (ICML)*. (2018)

—— Cited on pages 29, 32.

JEAN-YVES AUDIBERT and OLIVIER CATONI. Robust linear least squares regression. *The Annals of Statistics*. (2011). URL: <https://doi.org/10.1214/11-AOS918>

—— Cited on page 32.

ARINDAM BANERJEE. On Bayesian Bounds. *Proceedings of the 23rd international conference on Machine learning*. (2006)

—— Cited on page 25.

PETER BARTLETT and SHAHAR MENDELSON. Rademacher and Gaussian Complexities: Risk Bounds and Structural Results. *Conference on Computational Learning Theory (COLT)*. (2001)

—— Cited on page 22.

PETER BARTLETT and SHAHAR MENDELSON. Rademacher and Gaussian Complexities: Risk Bounds and Structural Results. *Journal of Machine Learning Research*. (2002)

—— Cited on page 22.

PETER L BARTLETT and WOLFGANG MAASS. Vapnik-Chervonenkis dimension of neural nets. *The handbook of brain theory and neural networks*. (2003)

—— Cited on page 23.

FELIX BIGGS and BENJAMIN GUEDJ. Differentiable PAC-Bayes Objectives with Partially Aggregated Neural Networks. *Entropy*. (2021). URL: <https://doi.org/10.3390/e23101280>

—— Cited on page 29.

FELIX BIGGS and BENJAMIN GUEDJ. Non-Vacuous Generalisation Bounds for Shallow Neural Networks. *Proceedings of the 39th International Conference on Machine Learning. PMLR*. (2022a). URL: <https://proceedings.mlr.press/v162/biggs22a.html>

—— Cited on page 29.

FELIX BIGGS and BENJAMIN GUEDJ. On Margins and Derandomisation in PAC-Bayes. *Proceedings of The 25th International Conference on Artificial Intelligence and Statistics*. (2022b). URL: <https://proceedings.mlr.press/v151/biggs22a.html>

—— Cited on page 29.

FELIX BIGGS and BENJAMIN GUEDJ. Tighter PAC-Bayes Generalisation Bounds by Leveraging Example Difficulty. *Proceedings of The 26th International Conference on Artificial Intelligence and Statistics. PMLR*. (2023). URL: <https://proceedings.mlr.press/v206/biggs23a.html>

—— Cited on page 29.

OLIVIER BOUSQUET and ANDRE ELISSEEFF. Algorithmic Stability and Generalization Performance. *Advances in Neural Information Processing Systems 13, Papers from Neural Information Processing Systems (NIPS) 2000, Denver, CO, USA*. (2000). URL: <https://proceedings.neurips.cc/paper/2000/hash/49ad23d1ec9fa4bd8d77d02681df5cfa-Abstract.html>

—— Cited on page 22.

References

OLIVIER CATONI. A PAC-Bayesian approach to adaptive classification. *preprint*. 840. (2003)

—— Cited on pages 23, 27.

OLIVIER CATONI. Statistical learning theory and stochastic optimization: Ecole d'Été de Probabilités de Saint-Flour, XXXI-2001. Vol. 1851. *Springer Science & Business Media*. (2004)

—— Cited on page 26.

OLIVIER CATONI. PAC-Bayesian supervised classification: the thermodynamics of statistical learning. *Institute of Mathematical Statistics*. (2007)

—— Cited on pages 23–25, 28, 32.

OLIVIER CATONI. PAC-Bayesian bounds for the Gram matrix and least squares regression with a random design. *CoRR*. abs/1603.05229. (2016)

—— Cited on page 32.

OLIVIER CATONI and ILARIA GIULINI. Dimension-free PAC-Bayesian bounds for matrices, vectors, and linear least squares regression. *CoRR*. abs/1712.02747. (2017)

—— Cited on page 32.

EUGENIO CLERICO, TYLER FARGHLY, GEORGE DELIGIANNIDIS, BENJAMIN GUEDJ, and ARNAUD DOUCET. Generalisation under gradient descent via deterministic PAC-Bayes. *arXiv preprint arXiv:2209.02525*. (2022)

—— Cited on page 30.

THOMAS M. COVER and JOY A. THOMAS. Elements of Information Theory. *Wiley*. (2001)

—— Cited on page 22.

IMRE CSISZÁR. I -Divergence Geometry of Probability Distributions and Minimization Problems. *The Annals of Probability*. (1975)

—— Cited on page 25.

NAN DING, XI CHEN, TOMER LEVINBOIM, SEBASTIAN GOODMAN, and RADU SORICUT. Bridging the Gap Between Practice and PAC-Bayes Theory in Few-Shot Meta-Learning. *Advances in Neural Information Processing Systems (NeurIPS)*. (2021)

—— Cited on pages 29, 32.

M. D. DONSKER and S. R. S. VARADHAN. Asymptotic evaluation of certain Markov process expectations for large time—III. *Communications on Pure and Applied Mathematics*. (1976)

——— Cited on page 25.

BENJAMIN DUPUIS and Umut ŞİMŞEKLI. Generalization Bounds for Heavy-Tailed SDEs through the Fractional Fokker-Planck Equation. *arXiv preprint arXiv:2402.07723*. (2024)

——— Cited on page 25.

RICK DURRETT. Probability: theory and examples. Vol. 49. *Cambridge university press*. (2019)

——— Cited on page 46.

GINTARE KAROLINA DZIUGAITE, ALEXANDRE DROUIN, BRADY NEAL, NITARSHAN RAJKUMAR, ETHAN CABALLERO, LINBO WANG, IOANNIS MITLIAGKAS, and DANIEL M. ROY. In search of robust measures of generalization. *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*. (2020). URL: <https://proceedings.neurips.cc/paper/2020/hash/86d7c8a08b4aaa1bc7c599473f5ddda-Abstract.html>

——— Cited on page 29.

GINTARE KAROLINA DZIUGAITE, KYLE HSU, WASEEM GHARBIH, GABRIEL ARPINO, and DANIEL ROY. On the role of data in PAC-Bayes bounds. *International Conference on Artificial Intelligence and Statistics (AISTATS)*. (2021)

——— Cited on page 29.

GINTARE KAROLINA DZIUGAITE and DANIEL ROY. Computing Nonvacuous Generalization Bounds for Deep (Stochastic) Neural Networks with Many More Parameters than Training Data. *Conference on Uncertainty in Artificial Intelligence (UAI)*. (2017)

——— Cited on pages 28–30, 32.

GINTARE KAROLINA DZIUGAITE and DANIEL ROY. Entropy-SGD optimizes the prior of a PAC-Bayes bound: Generalization properties of Entropy-SGD and data-dependent priors. *Proceedings of the 35th International Conference on Machine Learning*. (2018a). URL: <https://proceedings.mlr.press/v80/dziugaite18a.html>

——— Cited on page 30.

References

GINTARE KAROLINA DZIUGAITE and DANIEL M ROY. Data-dependent PAC-Bayes priors via differential privacy. *Advances in Neural Information Processing Systems. Curran Associates, Inc.* (2018b). URL: <https://proceedings.neurips.cc/paper/2018/file/9a0ee0a9e7a42d2d69b8f86b3a0756b1-Paper.pdf>

—— Cited on pages 32, 33.

MAHDI MILANI FARD and JOELLE PINEAU. PAC-Bayesian Model Selection for Reinforcement Learning. *Advances in Neural Information Processing Systems (NIPS)*. (2010)

—— Cited on pages 29, 31.

ALEC FARID and ANIRUDHA MAJUMDAR. Generalization Bounds for Meta-Learning via PAC-Bayes and Uniform Stability. *Advances in Neural Information Processing Systems (NeurIPS)*. (2021)

—— Cited on pages 29, 32.

HAMISH FLYNN, DAVID REEB, MELIH KANDEMIR, and JAN PETERS. PAC-Bayesian lifelong learning for multi-armed bandits. *Data Min. Knowl. Discov.* (2022). URL: <https://doi.org/10.1007/s10618-022-00825-4>

—— Cited on page 31.

PASCAL GERMAIN, ALEXANDRE LACASSE, FRANÇOIS LAVIOLETTE, and MARIO MARCHAND. PAC-Bayesian learning of linear classifiers. *International Conference on Machine Learning (ICML)*. (2009)

—— Cited on page 26.

PETER GRUNWALD, THOMAS STEINKE, and LYDIA ZAKYNTHINO. PAC-Bayes, MAC-Bayes and Conditional Mutual Information: Fast rate bounds that handle general VC classes. *Proceedings of Thirty Fourth Conference on Learning Theory*. Ed. by MIKHAIL BELKIN and SAMORY KPOTUFE. Vol. 134. *Proceedings of Machine Learning Research. PMLR*. (15–19 Aug 2021). URL: <https://proceedings.mlr.press/v134/grunwald21a.html>

—— Cited on page 22.

BENJAMIN GUEDJ. A Primer on PAC-Bayesian Learning. *Proceedings of the second congress of the French Mathematical Society*. (2019)

—— Cited on pages 23–25.

BENJAMIN GUEDJ and SYLVAIN ROBBIANO. PAC-Bayesian high dimensional bipartite ranking. *Journal of Statistical Planning and Inference*. 196. (2018). ISSN: 0378-

3758. DOI: <https://doi.org/10.1016/j.jspi.2017.10.010>. URL: <http://www.sciencedirect.com/science/article/pii/S0378375817301945>

—— Cited on page 26.

MERT GÜRBÜZBALABAN, Umut ŞİMŞEKLI, and LINGJIONG ZHU. The heavy-tail phenomenon in SGD. *International Conference on Machine Learning (ICML)*. (2021)

—— Cited on page 30.

MAXIME HADDOUCHE, BENJAMIN GUEDJ, OMAR RIVASPLATA, and JOHN SHAWE-TAYLOR. PAC-Bayes Unleashed: Generalisation Bounds with Unbounded Losses. *Entropy*. 23. (2021)

—— Cited on page 32.

MAHDI HAGHIFAM, BORJA RODRÍGUEZ-GÁLVEZ, RAGNAR THOBABEN, MIKAEL SKOGLUND, DANIEL M. ROY, and GINTARE KAROLINA DZIUGAITE. Limitations of Information-Theoretic Generalization Bounds for Gradient Descent Methods in Stochastic Convex Optimization. *Proceedings of The 34th International Conference on Algorithmic Learning Theory*. (2023). URL: <https://proceedings.mlr.press/v201/haghifam23a.html>

—— Cited on page 30.

URI HASSON, SAMUEL A NASTASE, and ARIEL GOLDSTEIN. Direct fit to nature: an evolutionary perspective on biological and artificial neural networks. *Neuron*. (2020)

—— Cited on pages 15, 19.

FREDRIK HELLSTRÖM, GIUSEPPE DURISI, BENJAMIN GUEDJ, and MAXIM RAGINSKY. Generalization bounds: Perspectives from information theory and PAC-Bayes. *arXiv preprint arXiv:2309.04381*. (2023)

—— Cited on page 23.

MATTHEW HOLLAND. PAC-Bayes under potentially heavy tails. *Advances in Neural Information Processing Systems (NeurIPS)* 32. Ed. by H. WALLACH, H. LAROCHELLE, A. BEYGELZIMER, F. D ALCHÉ-BUC, E. FOX, and R. GARNETT. *Curran Associates, Inc.* (2019). URL: <http://papers.nips.cc/paper/8539-pac-bayes-under-potentially-heavy-tails.pdf>

—— Cited on pages 26, 32.

GARETH JAMES, DANIELA WITTEN, TREVOR HASTIE, ROBERT TIBSHIRANI, *et al.* An introduction to statistical learning. Vol. 112. *Springer*. (2013)

—— Cited on page 21.

References

YIDING JIANG, BEHNAM NEYSHABUR, HOSSEIN MOBAHI, DILIP KRISHNAN, and SAMY BENGIO. Fantastic Generalization Measures and Where to Find Them. *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. (2020). URL: <https://openreview.net/forum?id=SJgIPJBFvH>

—— Cited on page 29.

ARYEH KONTOROVICH. Concentration in unbounded metric spaces and algorithmic stability. *Proceedings of the 31th International Conference on Machine Learning, ICML 2014, Beijing, China, 21-26 June 2014*. Vol. 32. *JMLR Workshop and Conference Proceedings*. JMLR.org. (2014). URL: <http://proceedings.mlr.press/v32/kontorovich14.html>

—— Cited on page 22.

ILJA KUZBORSKIJ and CSABA SZEPESVÁRI. Efron-Stein PAC-Bayesian Inequalities. *arXiv*. abs/1909.01931. (2019)

—— Cited on page 32.

Gael LETARTE, PASCAL GERMAIN, BENJAMIN GUEDJ, and FRANÇOIS LAVIOLETTE. Dichotomize and Generalize: PAC-Bayesian Binary Activated Deep Neural Networks. *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*. Ed. by HANNA M. WALLACH, HUGO LAROCHELLE, ALINA BEYGEZIMMER, FLORENCE D’ALCHÉ-BUC, EMILY B. FOX, and ROMAN GARNETT. (2019). URL: <https://proceedings.neurips.cc/paper/2019/hash/7ec3b3cf674f4f1d23e9d30c89426cce-Abstract.html>

—— Cited on page 29.

GUY LEVER, FRANÇOIS LAVIOLETTE, and JOHN SHAWE-TAYLOR. Distribution-dependent PAC-Bayes priors. *International Conference on Algorithmic Learning Theory*. Springer. (2010)

—— Cited on page 32.

GUY LEVER, FRANÇOIS LAVIOLETTE, and JOHN SHAWE-TAYLOR. Tighter PAC-Bayes bounds through distribution-dependent priors. *Theoretical Computer Science*. 473. (2013)

—— Cited on page 32.

BEN LONDON. A PAC-Bayesian Analysis of Randomized Learning with Application to Stochastic Gradient Descent. *Advances in Neural Information Processing Systems 30*:

Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA. (2017). URL: <https://proceedings.neurips.cc/paper/2017/hash/7fea637fd6d02b8f0adf6f7dc36aed93-Abstract.html>

—— Cited on page 30.

ZHOU LU, HONGMING PU, FEICHENG WANG, ZHIQIANG HU, and LIWEI WANG. The Expressive Power of Neural Networks: A View from the Width. Ed. by ISABELLE GUYON, ULRIKE VON LUXBURG, SAMY BENGIO, HANNA M. WALLACH, ROB FERGUS, S. V. N. VISHWANATHAN, and ROMAN GARNETT. (2017). URL: <https://proceedings.neurips.cc/paper/2017/hash/32cbf687880eb1674a07bf717761dd3a-Abstract.html>

—— Cited on pages 15, 19.

ANDREAS MAURER. A note on the PAC-Bayesian theorem. *arXiv. cs/0411099*. (2004)

—— Cited on pages 24, 26.

DAVID A MCALLESTER. Some PAC-Bayesian theorems. *Proceedings of the eleventh annual conference on Computational Learning Theory*. ACM. (1998)

—— Cited on page 23.

DAVID A MCALLESTER. PAC-Bayesian model averaging. *Proceedings of the twelfth annual conference on Computational Learning Theory*. ACM. (1999)

—— Cited on page 23.

DAVID A MCALLESTER. PAC-Bayesian Stochastic Model Selection. *Machine Learning*. (2003)

—— Cited on pages 23, 24.

ZAKARIA MHAMMEDI, PETER GRÜN WALD, and BENJAMIN GUEDJ. PAC-Bayes Un-Expected Bernstein Inequality. *Advances in Neural Information Processing Systems (NeurIPS)* 32. (2019). URL: <http://papers.nips.cc/paper/9387-pac-bayes-un-expected-bernstein-inequality.pdf>

—— Cited on pages 26, 32.

RADFORD M. NEAL. Bayesian learning for neural networks. *Springer Science & Business Media*. (2012)

—— Cited on page 22.

GERGELY NEU, GINTARE KAROLINA DZIUGAITE, MAHDI HAGHIFAM, and DANIEL M. ROY. Information-Theoretic Generalization Bounds for Stochastic Gradient De-

References

scent. *Proceedings of Thirty Fourth Conference on Learning Theory*. (2021). URL: <https://proceedings.mlr.press/v134/neu21a.html>

—— Cited on page 30.

BEHNAMEH NEYSHABUR, SRINADH BHOJANAPALLI, DAVID MCALLESTER, and NATI SREBRO. Exploring Generalization in Deep Learning. *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*. (2017). URL: <https://proceedings.neurips.cc/paper/2017/hash/10ce03a1ed01077e3e289f3e53c72813-Abstract.html>

—— Cited on page 29.

YUKI OHNISHI and JEAN HONORIO. Novel Change of Measure Inequalities with Applications to PAC-Bayesian Bounds and Monte Carlo Estimation. *International Conference on Artificial Intelligence and Statistics (AISTATS)*. (2021)

—— Cited on page 32.

LUCA ONETO, DAVIDE ANGUITA, and SANDRO RIDELLA. PAC-Bayesian analysis of distribution dependent priors: Tighter risk bounds and stability analysis. *Pattern Recognition Letters*. (2016)

—— Cited on page 32.

SEJUN PARK, CHULHEE YUN, JAEHO LEE, and JINWOO SHIN. Minimum Width for Universal Approximation. *9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3-7, 2021*. (2021). URL: <https://openreview.net/forum?id=0-XJwyoIF-k>

—— Cited on pages 15, 19.

EMILIO PARRADO-HERNÁNDEZ, AMIRAN AMBROLADZE, JOHN SHAWE-TAYLOR, and SHILIANG SUN. PAC-bayes bounds with data dependent priors. *Journal of Machine Learning Research*. (2012)

—— Cited on page 32.

ANASTASIA PENTINA and CHRISTOPH H. LAMPERT. A PAC-Bayesian bound for Lifelong Learning. *Proceedings of the 31th International Conference on Machine Learning, ICML 2014, Beijing, China, 21-26 June 2014*. Vol. 32. *JMLR Workshop and Conference Proceedings*. JMLR.org. (2014). URL: <http://proceedings.mlr.press/v32/pentina14.html>

—— Cited on page 31.

MARIA PEREZ-ORTIZ, OMAR RIVASPLATA, BENJAMIN GUEDJ, MATTHEW GLEESON, JINGYU ZHANG, JOHN SHAWE-TAYLOR, MIROSLAW BOBER, and JOSEF KITTLER. Learning PAC-Bayes Priors for Probabilistic Neural Networks. (2021a)

—— Cited on pages 29, 32.

MARIA PEREZ-ORTIZ, OMAR RIVASPLATA, EMILIO PARRADO-HERNANDEZ, BENJAMIN GUEDJ, and JOHN SHAWE-TAYLOR. Progress in Self-Certified Neural Networks. *NeurIPS 2021 Workshop on Bayesian Deep Learning*. (2021b)

—— Cited on page 29.

MARIA PEREZ-ORTIZ, OMAR RIVASPLATA, JOHN SHAWE-TAYLOR, and CSABA SZEPESVARI. Tighter Risk Certificates for Neural Networks. *Journal of Machine Learning Research*. (2021)

—— Cited on pages 28, 29.

ANTOINE PICARD-WEIBEL and BENJAMIN GUEDJ. On change of measure inequalities for f -divergences. *arXiv*. abs/2202.05568. (2022)

—— Cited on page 32.

FRIGYES RIESZ. *Leçons d'Analyse Fonctionnelle*. (1955)

—— Cited on pages 15, 19.

OMAR RIVASPLATA, ILJA KUZBORSKIJ, CSABA SZEPESVÁRI, and JOHN SHAWE-TAYLOR. PAC-Bayes analysis beyond the usual bounds. *Advances in Neural Information Processing Systems (NeurIPS)*. (2020)

—— Cited on page 31.

OMAR RIVASPLATA, VIKRAM M. TANKASALI, and CSABA SZEPESVARI. PAC-Bayes with Backprop. *arXiv*. (2019)

—— Cited on page 29.

BORJA RODRIGUEZ-GALVEZ, GERMAN BASSI, RAGNAR THOBABEN, and MIKAEL SKOGLUND. Tighter Expected Generalization Error Bounds via Wasserstein Distance. *Advances in Neural Information Processing Systems 34: Annual Conference on Neural Information Processing Systems 2021, NeurIPS 2021, December 6-14, 2021, virtual*. (2021)

—— Cited on page 22.

JONAS ROTHFUSS, VINCENT FORTUIN, MARTIN JOSIFOSKI, and ANDREAS KRAUSE. PACOH: Bayes-optimal meta-learning with PAC-guarantees. *International Conference*

References

on *Machine Learning (ICML)*. (2021)

—— Cited on pages 29, 32.

JONAS ROTHFUSS, MARTIN JOSIFOSKI, VINCENT FORTUIN, and ANDREAS KRAUSE. PAC-Bayesian Meta-Learning: From Theory to Practice. *arXiv*. abs/2211.07206. (2022)

—— Cited on pages 29, 32.

OTMANE SAKHI, PIERRE ALQUIER, and NICOLAS CHOPIN. PAC-Bayesian Offline Contextual Bandits With Guarantees. *International Conference on Machine Learning (ICML)*. (2023)

—— Cited on page 29.

YEVGENY SELDIN, NICOLÒ CESA-BIANCHI, PETER AUER, FRANÇOIS LAVIOLETTE, and JOHN SHAWE-TAYLOR. PAC-Bayes-Bernstein Inequality for Martingales and its Application to Multiarmed Bandits. *Proceedings of the Workshop on On-line Trading of Exploration and Exploitation 2*. PMLR. (2012a). URL: <https://proceedings.mlr.press/v26/seldin12a.html>

—— Cited on page 31.

YEVGENY SELDIN, FRANÇOIS LAVIOLETTE, NICOLÒ CESA-BIANCHI, JOHN SHAWE-TAYLOR, and PETER AUER. PAC-Bayesian Inequalities for Martingales. *IEEE Transactions on Information Theory*. (2012)

—— Cited on pages 29, 31.

YEVGENY SELDIN, FRANÇOIS LAVIOLETTE, JOHN SHAWE-TAYLOR, JAN PETERS, and PETER AUER. PAC-Bayesian Analysis of Martingales and Multiarmed Bandits. *arXiv*. abs/1105.2416. (2011)

—— Cited on pages 29, 31.

J. SHAWE-TAYLOR and R. C. WILLIAMSON. A PAC analysis of a Bayes estimator. *Proceedings of the 10th annual conference on Computational Learning Theory*. ACM. (1997)

—— Cited on page 23.

UMUT ŞİMŞEKLI, LEVENT SAGUN, and MERT GÜRBÜZBALABAN. A Tail-Index Analysis of Stochastic Gradient Noise in Deep Neural Networks. *International Conference on Machine Learning (ICML)*. (2019)

—— Cited on page 30.

NIKLAS THIEMANN, CHRISTIAN IGEL, OLIVIER WINTENBERGER, and YEVGENY SELDIN. A strongly quasiconvex PAC-Bayesian bound. *International Conference on Algorithmic Learning Theory*. PMLR. (2017)

—— Cited on page 26.

ILYA O. TOLSTIKHIN and YEVGENY SELDIN. PAC-Bayes-Empirical-Bernstein Inequality. *Advances in Neural Information Processing Systems (NeurIPS)*. (2013)

—— Cited on page 26.

VLADIMIR VAPNIK. An overview of statistical learning theory. *IEEE Trans. Neural Networks*. (1999). URL: <https://doi.org/10.1109/72.788640>

—— Cited on page 21.

VLADIMIR NAUMOVICH VAPNIK. The Nature of Statistical Learning Theory, Second Edition. *Statistics for Engineering and Information Science*. Springer. (2000)

—— Cited on page 22.

PAUL VIALARD, PASCAL GERMAIN, AMAURY HABRARD, and EMILIE MORVANT. A general framework for the practical disintegration of PAC-Bayesian bounds. *Machine Learning*. (2023)

—— Cited on pages 31, 32.

HAO WANG, MARIO DIAZ, JOSE CZNDIDO SILVEIRA SANTOS FILHO, and FLAVIO P. CALMON. An Information-Theoretic View of Generalization via Wasserstein Distance. *IEEE*. (2019). URL: <https://doi.org/10.1109/ISIT.2019.8849359>

—— Cited on page 22.

YI-SHAN WU and YEVGENY SELDIN. Split-kl and PAC-Bayes-split-kl Inequalities for Ternary Random Variables. *Advances in Neural Information Processing Systems*. (2022). URL: https://proceedings.neurips.cc/paper_files/paper/2022/file/49ffa271264808cf500ea528ed8ec9b3-Paper-Conference.pdf

—— Cited on page 26.

JINGZHAO ZHANG, SAI PRANEETH KARIMIREDDY, ANDREAS VEIT, SEUNGYEON KIM, SASHANK J. REDDI, SANJIV KUMAR, and SUVRIT SRA. Why are Adaptive Methods Good for Attention Models? *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*. (2020). URL: <https://proceedings.neurips.cc/paper/2020/hash/b05b57f6add810d3b7490866d74c0053-Abstract.html>

—— Cited on page 30.

References

SIJIA ZHOU, YUNWEN LEI, and ATA KABAN. Toward Better PAC-Bayes Bounds for Uniformly Stable Algorithms. *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023, New Orleans, LA, USA, December 10 - 16, 2023*. (2023). URL: http://papers.nips.cc/paper_files/paper/2023/hash/5e8309c9ca683e11672e3dbcd4b87776-Abstract-Conference.html

—— Cited on page 30.

Abstract. In machine learning, a model is learned from data to solve a task automatically. In the supervised classification setting, the model aims to predict the label associated with an input. The model is learned using a limited number of examples, each consisting of an input and its associated label. However, the model's performance on the examples, computed by the empirical risk, does not necessarily reflect the performance on the task, which is represented by the true risk. Moreover, since it is not computable, the true risk is upper-bounded by a generalization bound that mainly depends on two quantities: the empirical risk and a complexity measure. One way to learn a model is to minimize a bound by a type of algorithm called self-bounding. PAC-Bayesian bounds are well suited to the derivation of this type of algorithm. In this context, the first contribution consists in developing self-bounding algorithms that minimize PAC-Bayesian bounds to learn majority votes. If these bounds are well adapted to majority votes, their use for other models becomes less natural. To overcome this difficulty, a second contribution focuses on the disintegrated PAC-Bayesian bounds that are natural for more general models. In this framework, we provide the first empirical study of these bounds. In a third contribution, we derive bounds that allow us to incorporate complexity measures defined by the user.

Keywords. Machine Learning, Generalization, PAC-Bayesian Bound, Disintegrated PAC-Bayesian Bound, Self-Bounding Algorithm, Majority Vote, Neural Network, Complexity Measure.

Résumé. En apprentissage automatique, un modèle est appris à partir de données pour résoudre une tâche de manière automatique. Dans le cadre de la classification supervisée, le modèle vise à prédire la classe associée à une entrée. Le modèle est appris à l'aide d'un nombre limité d'exemples, chacun étant constitué d'une entrée et de sa classe associée. Cependant, la performance du modèle sur les exemples, calculée par le risque empirique, ne reflète pas nécessairement la performance sur la tâche qui est représentée par le risque réel. De plus, n'étant pas calculable, le risque réel est majoré pour obtenir une borne en généralisation qui dépend principalement de deux quantités : le risque empirique et une mesure de complexité. Une façon d'apprendre un modèle est de minimiser une borne par un type d'algorithme appelé auto-certié (ou auto-limitatif). Les bornes PAC-Bayésiennes sont bien adaptées à la dérivation de ce type d'algorithmes. Dans ce contexte, la première contribution consiste à développer des algorithmes auto-certiés qui minimisent des bornes PAC-Bayésiennes pour apprendre des votes de majorité. Si ces bornes sont bien adaptées aux votes de majorité, leur utilisation pour d'autres modèles devient moins naturelle. Pour pallier cette difficulté, une seconde contribution se concentre sur les bornes PAC-Bayésiennes désintégrées qui sont naturelles pour des modèles plus généraux. Dans ce cadre, nous apportons la première étude empirique de ces bornes. Dans une troisième contribution, nous dérivons des bornes permettant d'incorporer des mesures de complexité pouvant être définies par l'utilisateur.

Mot-clés. Apprentissage Automatique, Généralisation, Borne PAC-Bayésienne, Borne PAC-Bayésienne Désintégrée, Algorithme Auto-certié, Algorithme Auto-limitatif, Vote de Majorité, Réseau de Neurones, Mesure de Complexité.